

Variance reduction techniques for stochastic optimization

Matthew W. Hoffman

University of Cambridge

February 12, 2015

Variance reduction in Monte Carlo methods

We're interested in computing

$$\mathbb{E}_{p(x)} [f(x)] \approx \underbrace{\frac{1}{N} \sum_{i=1}^N f(x^i)}_F \quad \text{where } x^i \sim p(x)$$

but F is now a random variable

Problem: F may have high variance

Solution: replace F with a new quantity F' with the same expectation, but lower variance

$$\begin{aligned}\mathbb{E}[F'] &= \mathbb{E}[F] = \mathbb{E}[f(x)], \\ \text{var}[F'] &\leq \text{var}[F].\end{aligned}$$

Control variates

Consider an additional function $\phi(x)$ whose expectation $\mu_\phi = \mathbb{E}[\phi(x)]$ we know. We can introduce this function and write

$$\mathbb{E}[f(x)] = \underbrace{\mathbb{E}[f(x) - \phi(x)]}_{\text{use Monte Carlo here}} + \underbrace{\mu_\phi}_{\text{we know this}}$$

Nothing ground-breaking, but what about the variance?

$$\text{var}[f(x) - \phi(x)] = \text{var}[f(x)] - 2 \text{cov}[f(x), \phi(x)] + \text{var}[\phi(x)]$$

i.e. we can get a reduction in variance if f and ϕ are **strongly correlated**

ϕ is our **control variate**—so-called because it allows us to control the variance of our estimate

A few observations:

- we want $\text{cov} [f(x), \phi(x)] > \frac{1}{2} \text{var} [\phi(x)]$
- the control variate which minimizes the variance is easy, $\phi(x) = f(x)$, but this assumes we already know the integral
- instead we will often use simple variates of the form $a\phi$ and **optimize a assuming we are given ϕ**

Scaling the control variate

Now lets multiply the control variate by a scalar a ,

$$f'(x) = f(x) - a(\phi(x) - \mu_\phi)$$
$$\text{var} [f'(x)] = \text{var} [f(x)] - 2a \text{cov} [f(x), \phi(x)] + a^2 \text{var} [\phi(x)]$$

we can easily see by taking its derivative that this is minimized by

$$a = \frac{\text{cov} [f(x), \phi(x)]}{\text{var} [\phi(x)]}.$$

Plugging this in and dividing by the original variance we get

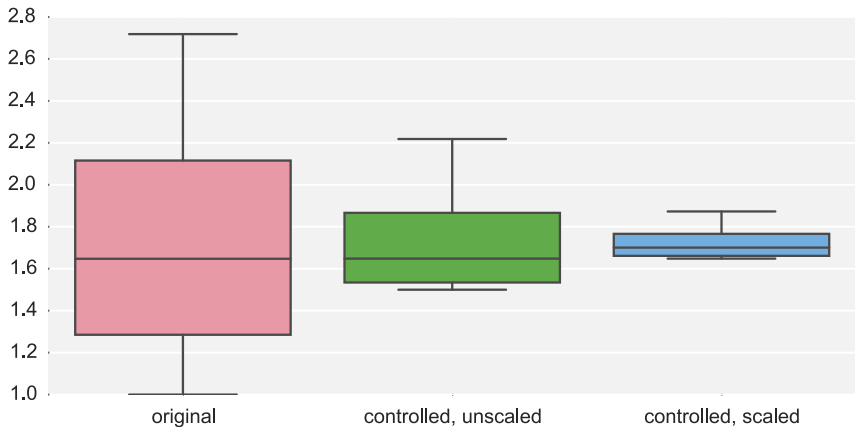
$$\frac{\text{var} [f'(x)]}{\text{var} [f(x)]} = 1 - \text{corr}^2 [f(x), \phi(x)]$$

i.e. the reduction in variance is directly related to their correlation.

- in computing a we typically won't have $\text{var}[\phi(x)]$ and $\text{cov}[f(x), \phi(x)]$ but we can still use their empirical estimates.

A simple example

Consider computing the expectation $\mathbb{E}[e^x]$ where $x \sim \mathcal{U}(0, 1)$ and use as control variate $\phi(x) = x - 0.5$.



Simple code generating the last plot:

```
x = np.random.rand(n)
f = np.exp(x)
phi = x - 0.5

# NOTE: requires touching all our data!
_, cov, _, var = np.cov([f, phi]).ravel()
f1 = f - phi
f2 = f - (cov/var) * phi
```

The comment is important; sometimes it may be too costly (or impossible) to view all our data.

Antithetic variables

as a special case of control variates

Consider a random variable given as a function of uniform variates

$$X = h(U_1, \dots, U_n), \text{ and}$$

$$W = h(1 - U_1, \dots, 1 - U_n)$$

We can use $0.5(X - W)$ as a control variate in order to estimate $\mathbb{E}[X]$

- also known as antithetic variables, resulting in the estimator $0.5(X + W)$
- this relies X and W being negatively correlated,
- provably reduces the variance when h is monotonic (either decreasing or increasing) in its inputs

Vector-valued variates

Given vector-valued $\mathbf{f}(x)$ we should use a control variate $\phi(x)$ with expectation μ_ϕ and an appropriately-sized matrix \mathbf{A} to define

$$\mathbf{f}'(x) = \mathbf{f}(x) - \mathbf{A}^\top(\phi(x) - \mu_\phi)$$

Let's say then that we select \mathbf{A} to minimize $\text{tr}[\text{cov}[\mathbf{f}'(x)]]$, leading to

$$\mathbf{A} = \Sigma^{-1}(\Omega + \Omega^\top)/2 \quad \text{where} \quad \begin{aligned} \Sigma &= \text{cov}[\phi(x)] \\ \Omega &= \text{cov}[\mathbf{f}(x), \phi(x)] \end{aligned}$$

This does require inversion of Σ though...

If we assume that \mathbf{A} is diagonal the optimal choice is given by

$$a_{ii} = \frac{\text{cov}[f_i(x), \phi_i(x)]}{\text{var}[\phi_i(x)]}$$

which is the same as the scalar case applied to each corresponding dimension

Assuming a single scalar value we get back

$$a = \frac{\sum_i \text{cov}[f_i(x), \phi_i(x)]}{\sum_i \text{var}[\phi_i(x)]}$$

which can be obtained by considering the scalar control variate case where we now just want to minimize the sum of variances (as done by Paisely et al.)

- for what follows I will assume $\mu_\phi = \mathbb{E}[\phi(x)] = 0$ and that the control variate contains any multiplier a

Stochastic gradient descent (SGD)

Often we find ourselves wanting to minimize some expected cost function

$$J(\theta) = \mathbb{E}_{p(x)} [c_{\theta}(x)]$$

whose gradient is given by

$$\begin{aligned}\nabla J(\theta) &= \int p(x) \nabla c_{\theta}(x) dx \\ &\approx \frac{1}{N} \sum_{i=1}^N \nabla c_{\theta}(x^i) \quad \text{for } x^i \sim p(\cdot)\end{aligned}$$

This is just a Monte Carlo estimate of something (that just happens to be a gradient)! So we can apply a suitable control variate.

SGD for logistic regression

For logistic regression we want to classify inputs \mathbf{x} as one of two classes, $y \in \{-1, 1\}$, the likelihood for which is

$$p_{\theta}(y|\mathbf{x}) = \sigma(y\boldsymbol{\theta}^T \mathbf{x}) \quad \sigma(z) = (1 + \exp(-z))^{-1} \text{ is the logistic}$$

The cost for a single observation is its log-likelihood, with gradient

$$\nabla c_{\theta}(\mathbf{x}, y) = y\mathbf{x}\sigma(-y\boldsymbol{\theta}^T \mathbf{x})$$

The noisy gradient is

$$\nabla J(\theta) \approx \frac{1}{N} \sum_{i=1}^N \nabla c_{\theta}(\mathbf{x}^i, y^i)$$

SGD for logistic regression

For logistic regression we want to classify inputs \mathbf{x} as one of two classes, $y \in \{-1, 1\}$, the likelihood for which is

$$p_{\theta}(y|\mathbf{x}) = \sigma(y\boldsymbol{\theta}^T \mathbf{x}) \quad \sigma(z) = (1 + \exp(-z))^{-1} \text{ is the logistic}$$
$$\hat{\sigma}(z) = \sigma(\hat{z})(1 + \sigma(-\hat{z})(z - \hat{z}))$$

is its 1st-order Taylor exp.

The cost for a single observation is its log-likelihood, with gradient

$$\nabla c_{\theta}(\mathbf{x}, y) = y\mathbf{x}\sigma(-y\boldsymbol{\theta}^T \mathbf{x})$$
$$\phi(\mathbf{x}, y) = y\mathbf{x}\hat{\sigma}(-y\boldsymbol{\theta}^T \mathbf{x}) - \mu_{\phi}$$

The noisy gradient is

$$\nabla J(\theta) \approx \frac{1}{N} \sum_{i=1}^N \nabla c_{\theta}(\mathbf{x}^i, y^i) - \phi(\mathbf{x}^i, y^i)$$

What did I leave out?

- the expectation μ_ϕ requires the mean and variance of both positive and negative inputs; **this requires a full pass over the data**
- note though that we need only do this once

SGD with parameterized distributions

Consider an objective where the distribution itself is parameterized

$$J(\theta) = \mathbb{E}_{p_{\theta}(x)} [c(x)]$$

and whose gradient is

$$\begin{aligned}\nabla J(\theta) &= \int \nabla p_{\theta}(x) c(x) dx \\ &= \int p_{\theta}(x) \nabla \log p_{\theta}(x) c(x) dx \\ &\approx \frac{1}{N} \sum_{i=1}^N \nabla \log p_{\theta}(x^i) c(x^i) \quad \text{for } x^i \sim p_{\theta}(\cdot)\end{aligned}$$

An aside: so far this has **nothing to do with control variates** or variance reduction in any way. We are just evaluating a gradient.

SGD for policy learning: REINFORCE, GPOMDP

We can now apply this to reinforcement learning where

- $x = (s_{0:T}, a_{0:T})$ represents a trajectory
- $c(x) = -\sum_t \gamma^t r(s_t, a_t)$ are summed (discounted) rewards
- the probability of trajectories has Markovian structure,

$$p_{\theta}(x) = \mu(s_0) \prod_t \pi_{\theta}(a_t | s_t) p(s_{t+1} | s_t, a_t)$$

Plugging this into the previous framework we get,

$$\nabla J(\theta) \approx -\frac{1}{N} \sum_{i=1}^N \left[\sum_t \nabla \log \pi_{\theta}(a_t^i | s_t^i) \right] \left[\sum_t \gamma^t r(s_t^i, a_t^i) \right]$$

this is **not quite** the REINFORCE algorithm

Eliminating expectations in REINFORCE

Let $z_k = (s_k, a_k)$ be a state/action pair at time k . The previous gradient is a sum of many “cross-time” terms,

$$\begin{aligned} \mathbb{E}[\gamma^k r(z_k) \nabla \log \pi_{\theta}(z_t)] \quad \text{for } k < t \\ = \int p_{\theta}(z_k) \gamma^k r(z_k) \underbrace{\left[\int p_{\theta}(z_t|z_k) \nabla \log \pi_{\theta}(z_t) dz_t \right]}_{\text{expectation of a score}} dz_k \end{aligned}$$

By eliminating these terms (their expectation is zero!) we get REINFORCE,

$$\nabla J(\theta) \approx -\frac{1}{N} \sum_{i=1}^N \sum_{t=0}^T \sum_{k=t}^T \nabla \log \pi_{\theta}(a_t^i | s_t^i) \gamma^k r(s_k^i, a_k^i)$$

Note: if we didn't eliminate these terms they would **only add variance**

Control variates in REINFORCE (baselines)

In the same way that we eliminated zero-mean terms in the previous slide we can also add terms,

$$\nabla J(\theta) \approx -\frac{1}{N} \sum_{i=1}^N \sum_{t=0}^T \sum_{k=t}^T \nabla \log \pi_{\theta}(a_t^i | s_t^i) \left[\gamma^k r(s_k^i, a_k^i) - \hat{b}_k(s_k^i, a_k^i) \right]$$

which is called a **baseline**, i.e. a “baseline reward” to improve on

This can be interpreted as a control variate of the form

$$\phi(x) = \sum_{t=0}^T \sum_{k=t}^T \nabla \log \pi_{\theta}(a_t | s_t) \hat{b}_k(s_k, a_k)$$

which so long as b_k is computed using only state/action pairs **before** time k will have expectation zero

Choice of baseline

There is some analysis in Greensmith et al. providing an **optimal baseline** under various settings—a bit complicated (and different from the earlier analysis)

However, a common baseline to use is the averaged reward:

$$\hat{b}_k = \frac{1}{N} \sum_{i=1}^N \sum_{t=0}^K \gamma^t r(s_t^i, a_t^i)$$

in some sense this is intuitive and gives rise to the **baseline** name:

by combining this with our gradient the reward provides us with an improvement over the average

Actor-critic methods

Another technique involves using the value function as a baseline,

$$V^{\pi}(s) = \mathbb{E} \left[\sum_{t=0}^{\infty} \gamma^t r(s_t, a_t) | s_0 = s \right]$$

which is similar to the averaged-reward baseline presented earlier

Actor-critic methods extend this to using compatible function approximation for the value-function (approximate using a linear function of the policy gradient)

The Natural Actor-Critic takes these ideas and applies the *natural gradient*. Whether this counts as a *variance reduction* technique is a bit murky.

References I

- J. Baxter and P. L. Bartlett. Infinite-horizon policy-gradient estimation. *Journal of Artificial Intelligence Research*, pages 319–350, 2001.
- V. R. Konda and J. N. Tsitsiklis. Actor-critic algorithms. In *Advances in Neural Information Processing Systems*, volume 13, pages 1008–1014, 2000.
- J. Paisley, D. Blei, and M. Jordan. Variational bayesian inference with stochastic search. In *the International Conference on Machine Learning*, 2012.
- J. Peters and S. Schaal. Policy gradient methods for robotics. In *International Conference on Intelligent Robots and Systems*, pages 2219–2225, 2006.
- S. M. Ross. *Simulation*. Academic Press, 4 edition, 2006.

References II

- R. S. Sutton, D. A. McAllester, S. P. Singh, Y. Mansour, et al. Policy gradient methods for reinforcement learning with function approximation. In *Advances in Neural Information Processing Systems*, volume 13, pages 1057–1063, 2000.
- C. Wang, X. Chen, A. J. Smola, and E. P. Xing. Variance reduction for stochastic gradient optimization. In *Advances in Neural Information Processing Systems*, pages 181–189, 2013.
- R. J. Williams. Simple statistical gradient-following algorithms for connectionist reinforcement learning. *Machine learning*, 8(3-4): 229–256, 1992.